



# Service Orchestration

## BigData from testing to production

PhD. Ronald Muresano

# Course Roadmap

**Orchestration for developing Docker** and Docker-Compose.

## Big Data Challenges

An overview applied to the SMEs. Orchestration Philosophy

## Docker Swarm

Orchestration in a real private cluster (Deploying and Creating services)

## Kubernetes

Production Orchestration. Automating deployment, scaling, and management of containerized applications

## Cloud Infrastructure

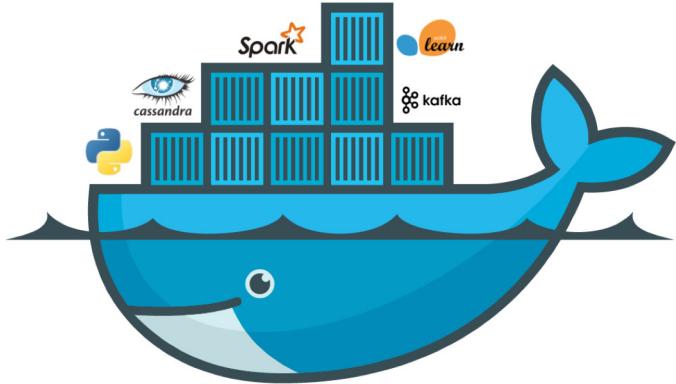
Google Cloud and AWS deployments



# Docker

It's a software platform which allows to packages application with its dependencies with the objective of deploying them easily.

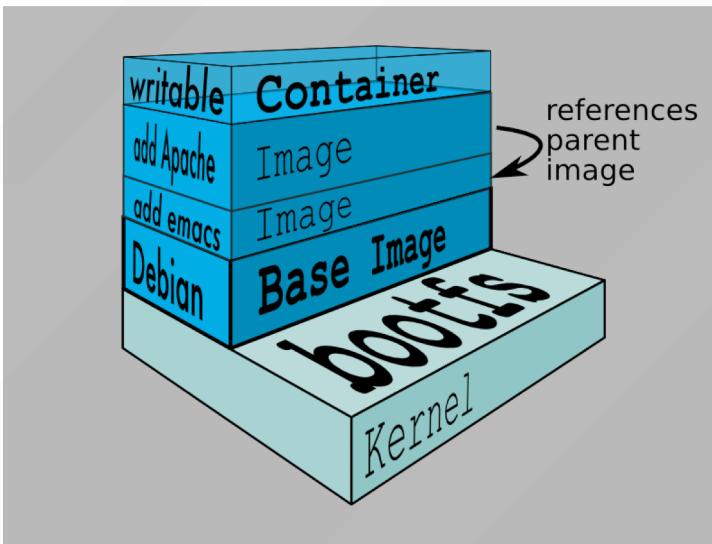
These packages are called containers.



- ❖ Portability of your compute environment
- ❖ Self-sufficient
- ❖ Light-Weight
- ❖ Flexible and scalable
- ❖ Strengthen your engineering steps
- ❖ Reproducibility ( As Data Scientists)

# Docker image

A docker image is a self-contained package in which is included all the necessary tools, code, runtime, libraries, environment variables, config files, dependencies, etc.



- ❖ An image is conformed by a set of layers which are Superimposition one to another.
- ❖ IMPORTANT: each instruction is a layer, in this sense, we can minimize the image size if some of the layer are grouped.
- ❖ A new layer is created when an image is modified.
- ❖ From an image, it can be created N containers that execute the image's content

**To inspect the docker image layers:** docker history <name of the image>

# Docker Advantages

## Standardization and Productivity

Ensure consistency across multiple development and release cycles, standardizing your environment.

## Rapid Deployment

Docker manages to reduce deployment to seconds.

## CI Efficiency

Enables you to build a container image and use that same image across every step of the deployment process

## Continuous Deployment and Testing

Ensures consistent environments from development to production

## Security

From a security point of view, it ensures that applications that are running on containers are completely segregated and isolated from each other

## Compatibility and Maintainability

Eliminate the “it works on my machine” problem once and for all. One of the benefits that the entire team will appreciate is parity.

## Multi-Cloud Platforms

One of Docker’s greatest benefits is portability.

## Scalable

It can increase and automatically distribute the container replication

## Simplicity and Faster Configurations

One of the key benefits of Docker is the way it simplifies matters

## Isolation

Docker ensures your applications and resources are isolated and segregated.

# Docker Architecture

It uses a client-server architecture, where the client communicates with the docker daemon in order to execute, build and distribute the containers.

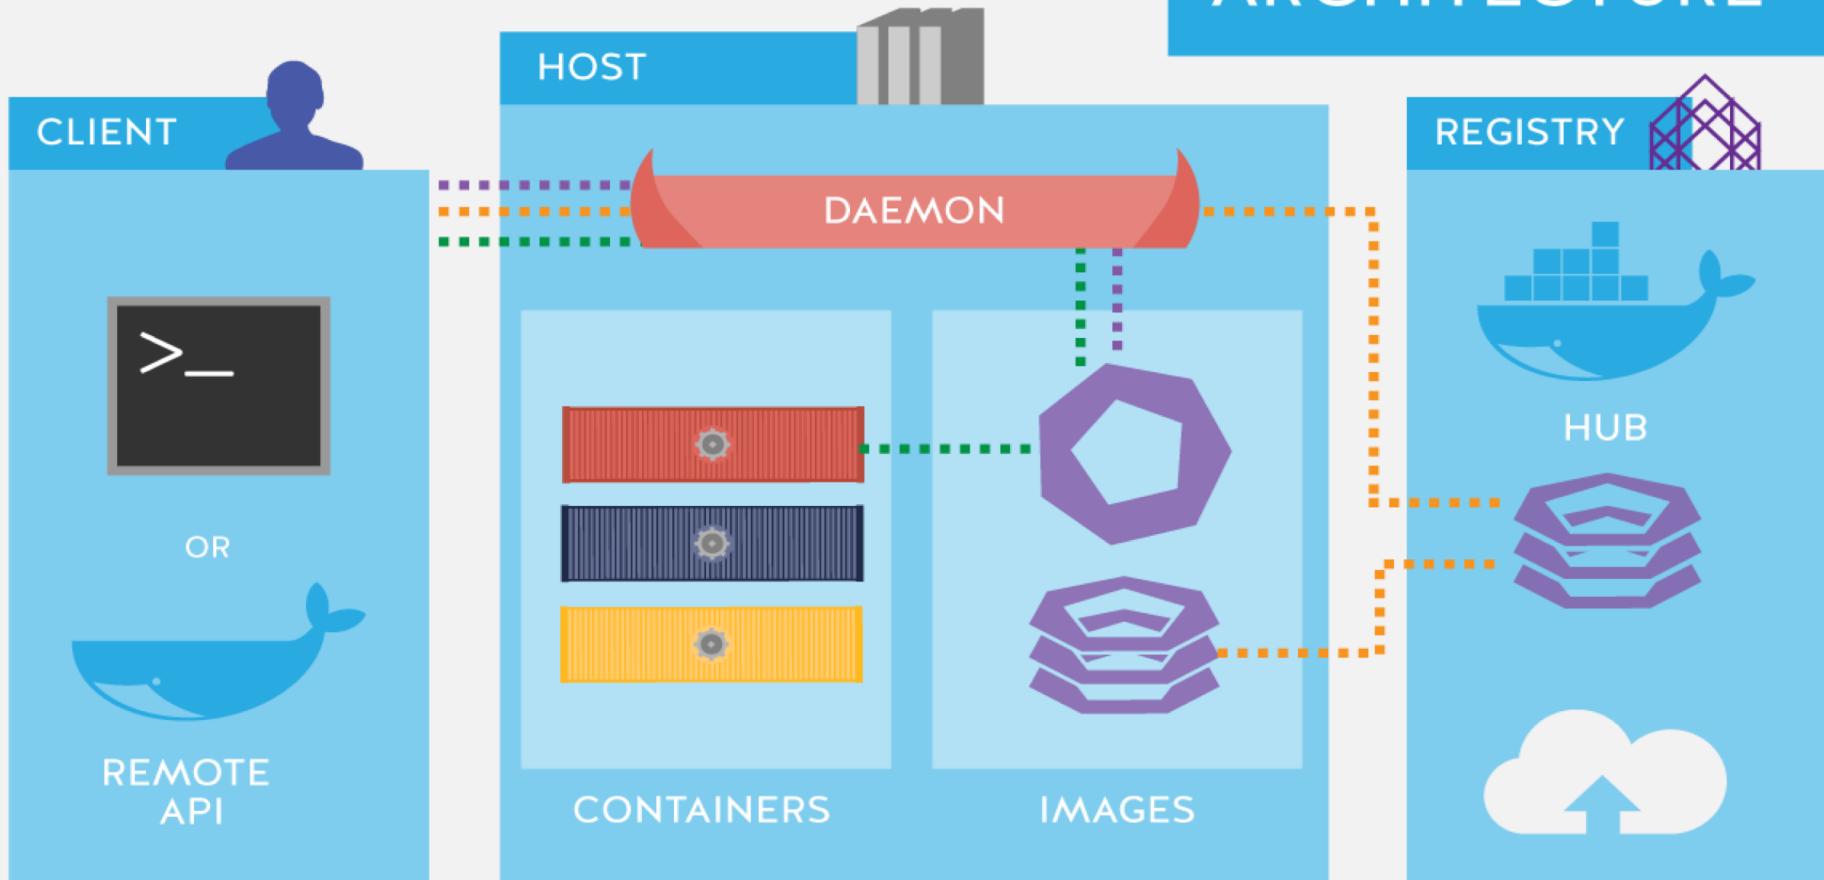
The communication between the client and the daemon is through an API-Rest. They use sockets UNIX, if they are running in the same system or TCP when a client talk with a remote daemon.

BUILD

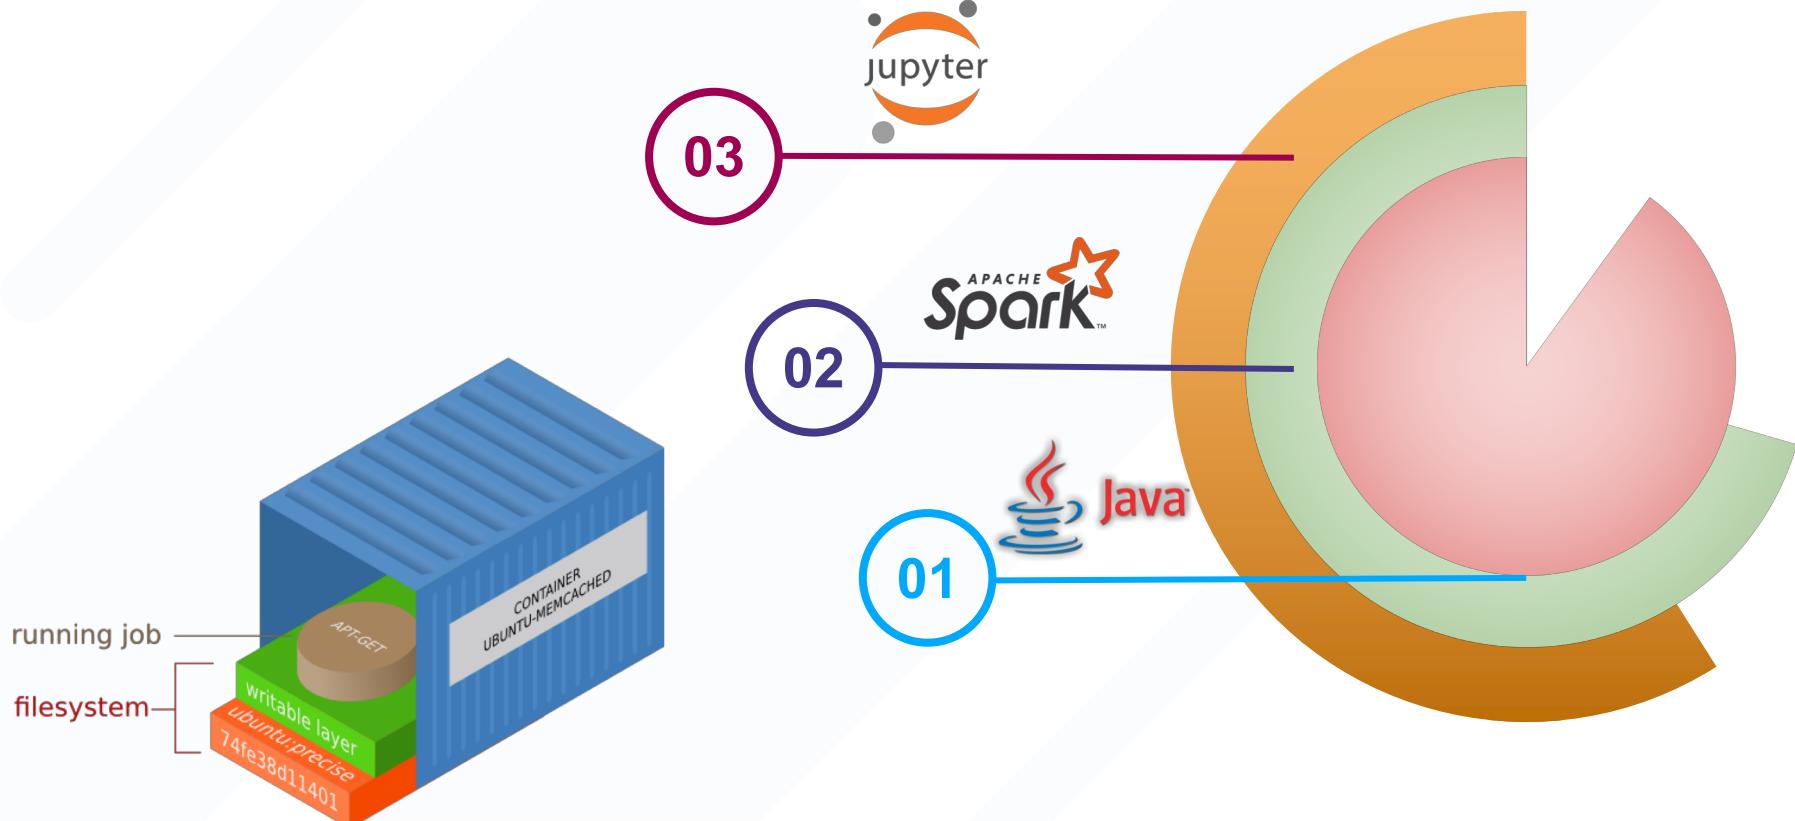
PULL

RUN

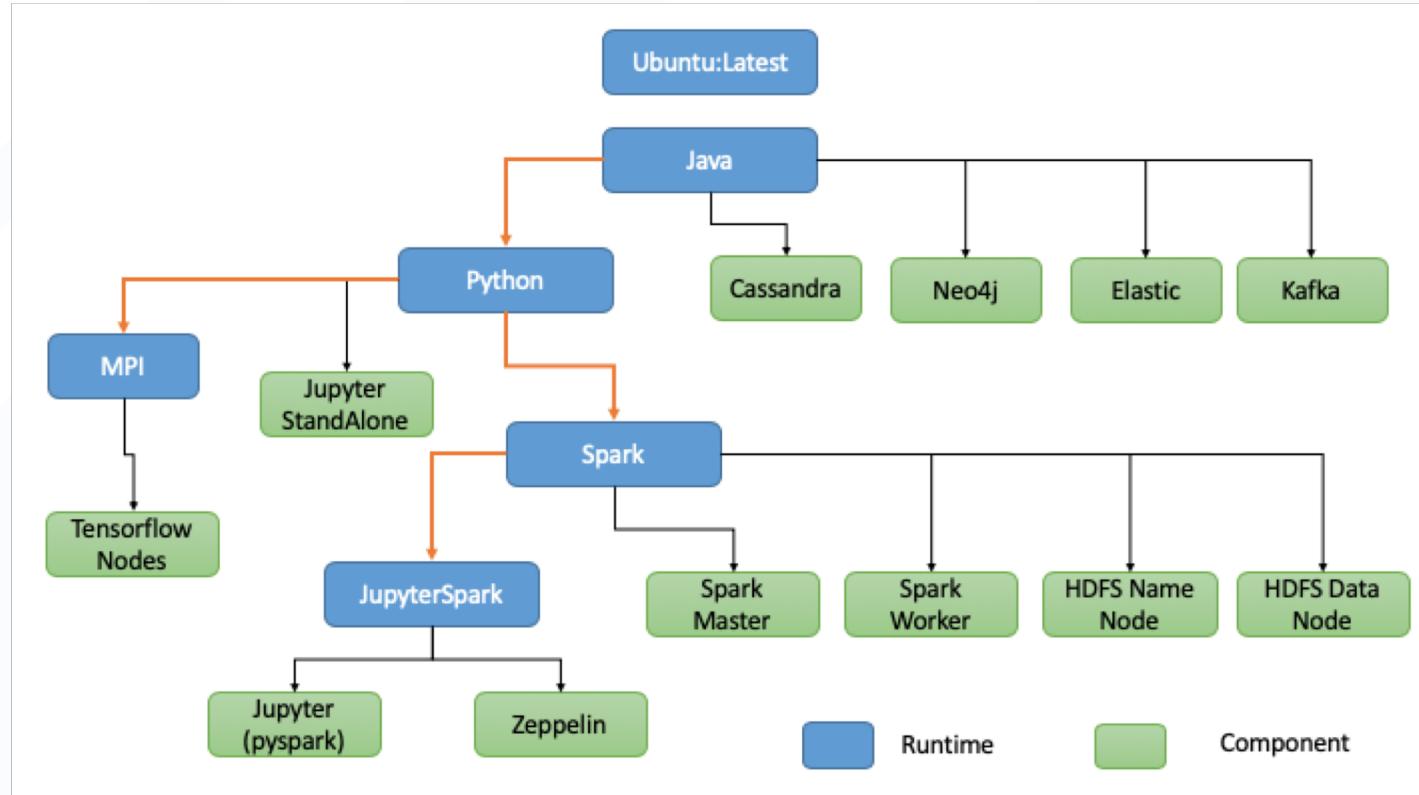
# DOCKER ARCHITECTURE



# Docker image (Runtime for BDaaS)



# Docker image (Runtime for BDaaS)



# Docker image (Runtime for BDaaS)

```
rmuresano@rmuresano-pc:~/git$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
rmuresano/bdaasjupyter    0_0_1   2ac679c04de1  3 weeks ago  2.73GB
rmuresano/bdaasspark     0_0_1   f301f6a0a6a5  3 weeks ago  1.54GB
rmuresano/bdaaspython     0_0_1   7947fff07802  3 weeks ago  1.02GB
rmuresano/bdaasjava      0_0_1   20ec23cf25bf  3 weeks ago  588MB
radiatus.iti.es/runtime/mpi 0_0_25  80958b89f09b  4 weeks ago  3.83GB
radiatus.iti.es/runtime/jupyterspark 0_0_25  fe8e0c03074a  4 weeks ago  3.11GB
radiatus.iti.es/runtime/spark    0_0_25  afffb73b17000 4 weeks ago  1.45GB
radiatus.iti.es/runtime/hadoop   0_0_25  8ab082177c47  4 weeks ago  1.28GB
radiatus.iti.es/runtime/python   0_0_25  7a35cb9205c2  4 weeks ago  998MB
radiatus.iti.es/runtime/mysql   0_0_25  a63fadab2f60  4 weeks ago  939MB
radiatus.iti.es/runtime/basic   0_0_25  d7e8b05df2d2  4 weeks ago  423MB
ubuntu                latest   7698f282e524  7 weeks ago  69.9MB
nginx                 latest   53f3fd8007f7  2 months ago  109MB
node                  lts-slim  8bfde27ed7db  2 months ago  143MB
saasdk/local-stamp      latest   4dec561fdf0f  4 months ago  848MB
dockersamples/visualizer  latest   f6411ebd974c  6 months ago  166MB
saasdk/forge.domains    latest   7dc75f09bb01  7 months ago  206MB
registry               2        2e2f252f3c88  9 months ago  33.3MB
saasdk/eslap.cloud_runtime native dev 2_0_0   383c9709f1ea  17 months ago  525MB
saasdk/eslap.cloud_runtime_native 2_0_0   71cea2ae9bbc  17 months ago  216MB
saasdk/eslap.cloud_elk      1_0_0   d6c67b531589  2 years ago   594MB
eslap.cloud/elk           1_0_0   d6c67b531589  2 years ago   594MB
rmuresano@rmuresano-pc:~/git$ ++■
```

# Docker image (Runtime for BDaaS)

Example 1:

```
rmuresano@rmuresano-pc:~/git$ docker run --rm -ti debian
Unable to find image 'debian:latest' locally
latest: Pulling from library/debian
6f2f362378c5: Pull complete
Digest: sha256:118cf8f3557elea766c02f36f05f6ac3e63628427ea8965fb861be904ec35a6f
Status: Downloaded newer image for debian:latest
root@c056b43ddffa:/#
```

```
root@c056b43ddffa:/# cat /etc/os-release
PRETTY_NAME="Debian GNU/Linux 9 (stretch)"
NAME="Debian GNU/Linux"
VERSION_ID="9"
VERSION="9 (stretch)"
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
root@c056b43ddffa:/#
```

# Docker image (Runtime for BDaaS)

debian                    latest                    e1de74e67cc7                    3 weeks ago                    101MB

A complete documentation can be found:

<https://docs.docker.com/engine/reference/commandline/cli/>

## Docker management

**Docker run:** Executes an order to run a new container.

**Docker exec:** Invokes an order in a running container.

**Docker ps:** Lists containers

**Docker rm:** Deletes container with an specific tag

**Docker rmi:** Deletes images. Sometimes must be forced.

**Docker images:** List Images (For us Runtimes)

**Docker build:** Build an image from a Dockerfile

## Docker Registry management

**Docker tag:** It control and manage the image tagging

**Docker login:** Identify the in the Docker registry

**Docker pull and push:** Push and Pull an image to/from the registry

**Docker search:** Search an image in the Docker registry

**Docker port:** List the associated ports to an specific container

**Docker network:** Manages the network interconnection between containers

**Docker volume:** Shows the volume information

**Docker inspect:** Returns the low level information about a docker container.

# Docker image (Runtime for BDaaS)

NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
elasticsearch	Elasticsearch is a powerful open source se...	3736	[OK]	
nshou/elasticsearch-kibana	Elasticsearch-7.1.1 Kibana-7.1.1	102		[OK]
itzg/elasticsearch	Provides an easily configurable Elasticsea...	67		[OK]
mobz/elasticsearch-head	elasticsearch-head front-end and standalon...	46		
elastichq/elasticsearch-hq	Official Docker image for ElasticHQ: Elast...	31		[OK]
lmenezes/elasticsearch-kopf	elasticsearch kopf	18		[OK]
bitnami/elasticsearch	Bitnami Docker Image for Elasticsearch	18		[OK]
elastic/elasticsearch	The Elasticsearch Docker image maintained ...	17		
barnybug/elasticsearch	Latest Elasticsearch 1.7.2 and previous re...	17		[OK]
taskrabbit/elasticsearch-dump	Import and export tools for elasticsearch	16		[OK]
esystemstech/elasticsearch	Debian based Elasticsearch packing for Lif...	15		
monsantoco/elasticsearch	ElasticSearch Docker image	11		[OK]
justwatch/elasticsearch_exporter	Elasticsearch stats exporter for Prometheus	9		
mesoscloud/elasticsearch	[UNMAINTAINED] Elasticsearch	9		[OK]
blacktop/elasticsearch	Alpine Linux based Elasticsearch Docker Image	8		[OK]
centerforopenscience/elasticsearch	Elasticsearch	4		[OK]
barchart/elasticsearch-aws	Elasticsearch AWS node	3		
phenompeople/elasticsearch	Elasticsearch is a powerful open source se...	1		[OK]
bitnami/elasticsearch-exporter	Bitnami Elasticsearch Exporter Docker Image	1		[OK]
jetstack/elasticsearch-pet	An elasticsearch image for kubernetes PetSets	1		[OK]
18fgsa/elasticsearch-ha	Built from https://github.com/18F/kubernetes	0		
18fgsa/elasticsearch	Built from https://github.com/docker-libra...	0		
wreulicke/elasticsearch	elasticsearch	0		[OK]
backplane/elasticsearch-curator	Elasticsearch Curator (https://github.com/...	0		
axway/elasticsearch-docker-beat	"Beat" extension to read logs of container...	0		[OK]

# Docker Management (Runtime for BDaaS)

## Example 2:

### Deployment Information

```
rmuresano@rmuresano-pc:~/git$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
 NAMES
972e8f4c42fc        rmuresano/bdaasjupyter:0_0_4   "/bin/sh -c ./star..."   6 seconds ago      Up 5 seconds       0.0.0.0:8900->8900/tcp
jupyter
2732b95e72f7        rmuresano/bdaassparkworker:0_0_4   "/bin/sh -c ./star..."   6 seconds ago      Up 5 seconds       0.0.0.0:8081->8081/tcp
jupyterspark_sparkworker_1
ad0f7269fd06        rmuresano/bdaassparkmaster:0_0_4   "/bin/sh -c ./star..."   6 seconds ago      Up 6 seconds       0.0.0.0:7077->7077/tcp, 0.0.0.0:8080->8080/tcp, 0.0.0.0:9000->9000/tcp, 0.0.0.0:50070->50070/tcp
sparkmaster
```

### Port

```
rmuresano@rmuresano-pc:~/git$ docker port 972e8f4c42fc
8900/tcp -> 0.0.0.0:8900
```

### Volume Information

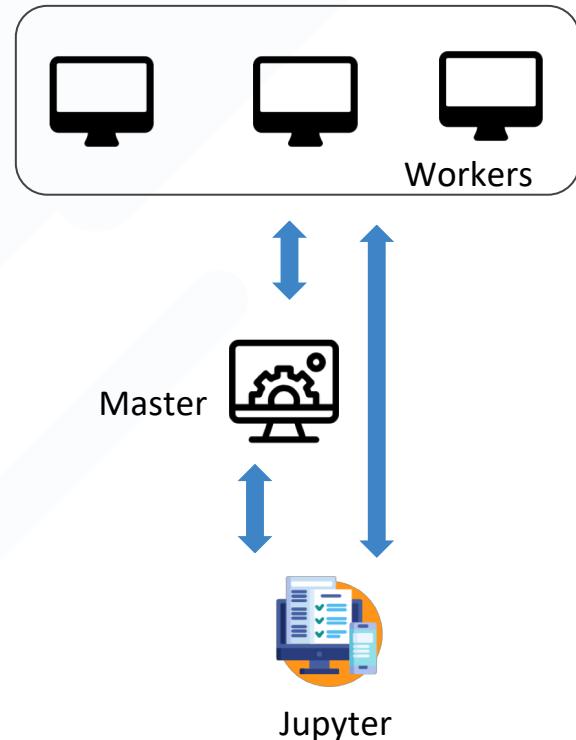
```
rmuresano@rmuresano-pc:~/git$ docker volume ls
DRIVER          VOLUME NAME
local           jupyterspark_hdfs-name-node
local           jupyterspark_jupyter-vol
```

```
rmuresano@rmuresano-pc:~/git$ docker volume inspect jupyterspark_jupyter-vol
[
    {
        "CreatedAt": "2019-06-05T15:32:58+02:00",
        "Driver": "local",
        "Labels": null,
        "Mountpoint": "/var/lib/docker/volumes/jupyterspark_jupyter-vol/_data",
        "Name": "jupyterspark_jupyter-vol",
        "Options": {},
        "Scope": "local"
    }
]
```

# Docker Management (Runtime for BDaaS)

```
rmuresano@rmuresano-pc:~$ docker network inspect 9b0d03c3bfaf
[{"Name": "jupyterspark_default",
 "Id": "9b0d03c3bfaf58593965866abdc0827c172f56e536facc2bab4d68a30b36982",
 "Created": "2019-06-05T14:39:05.3651268+02:00",
 "Scope": "local",
 "Driver": "bridge",
 "EnableIPv6": false,
 "IPAM": {
   "Driver": "default",
   "Options": null,
   "Config": [
     {
       "Subnet": "172.18.0.0/16",
       "Gateway": "172.18.0.1"
     }
   ],
   "Internal": false,
   "Attachable": true,
   "Ingress": false,
   "ConfigFrom": {
     "Network": ""
   },
   "ConfigOnly": false,
   "Containers": [
     "2732b95e72f7b96afbbb848576b7d25daee2205394ae87af515989f6bc53b92": {
       "Name": "jupyterspark_sparkworker_1",
       "EndpointID": "1fc590e352834f4325fb0a947628ec4cbadce45abe361eb58ff62a95eb540eb2",
       "MacAddress": "02:42:ac:12:00:03",
       "IPv4Address": "172.18.0.3/16",
       "IPv6Address": ""
     },
     "972e8f4c42fc34f0e386be00d9f71950203bde1be9b6c49f51ebc695871eca1": {
       "Name": "jupyter",
       "EndpointID": "c2201cfcfb198be7296c6232aef29b6e6d0ef7aafd67121b12b98e9123d345b",
       "MacAddress": "02:42:ac:12:00:04",
       "IPv4Address": "172.18.0.4/16",
       "IPv6Address": ""
     },
     "ad0f7269fd0de203ed96de519881ef8eae9bc1c1654cf8124068c7b71daf2295": {
       "Name": "sparkmaster",
       "EndpointID": "2645656d33b647b34796e585403e0b530054347b55ea9127d4a10826d3750415",
       "MacAddress": "02:42:ac:12:00:02",
       "IPv4Address": "172.18.0.2/16",
       "IPv6Address": ""
     }
   ],
   "Options": {},
   "Labels": {
     "com.docker.compose.network": "default",
     "com.docker.compose.project": "jupyterspark"
   }
  ]
}
```

Container  
network  
specification



# Docker Runtime building for BDaaS (DockerFile)

- To generate an image, we have to use the instruction docker build.
- The Dockerfile is a text document which contains a set of instructions in order to assemble an image.
- The context is the set of files that are in an specific route (local directory) o an URL (git repository reference)
- The context is processed in a recursive manner. This means if a directory is included all sub-directory are also included.

# Docker Runtime building for BDaaS (DockerFile)

- Like a Makefile (shell script with keywords)
- Extends from a Base Image
- Results in a new Docker Image
- Imperative, not Declarative
- Docker file lists the steps needed to build an images
- docker build is used to run a Docker file
- It can define default command for docker run, ports to expose, etc.

# Creating a Runtime from a DockerFile (1/2)

## Example 3:

Using the following invocation the docker client create a new tagged image with the name defined. For example, rmuresano/<name-of-the-image> and the version 0\_0\_1.

It is done passing the working directory as the current context and using the DockerFile.

```
muru@rmuresano-pc:~/git/bdaasmicroservices/docker/runtimes/helloExample$ docker build -t my_docker_flask:latest .
rmuresano@rmuresano-pc:~/git/bdaasmicroservices/docker/runtimes/helloExample$ docker build -t my_docker_flask:latest .
Sending build context to Docker daemon 5.12KB
Step 1/6 : FROM python
--> 34a518642c76
Step 2/6 : COPY . /app
--> cc32ab3aca88
Step 3/6 : WORKDIR /app
--> cfc472c244b1
Removing intermediate container d9bb2d8d7e92
Step 4/6 : RUN pip install -r requirements.txt
--> Running in c8960b77caa7
Collecting flask (from -r requirements.txt (line 1))
  Downloading https://files.pythonhosted.org/packages/c3/31/6904ac846fc65a7fa6cac8b4ddc392ce96ca08ee67b0f97854e9575bbb26/Flask-1.1.0-py2.py3-none-any.whl (94kB)
Collecting flask_restful (from -r requirements.txt (line 2))
  Downloading https://files.pythonhosted.org/packages/17/44/6e490156ee443ca81d5f88b61bb4bb133d44d75b0b716ebe92489508da4/Flask_RESTful-0.3.7-py2.py3-none-any.whl
Collecting Werkzeug>=0.15 (from flask->-r requirements.txt (line 1))
  Downloading https://files.pythonhosted.org/packages/9f/57/92a497e38161ce40606c27a86759c6b92dd34fcdb33f64171ec559257c02/Werkzeug-0.15.4-py2.py3-none-any.whl (327kB)
Collecting click>=5.1 (from flask->-r requirements.txt (line 1))
  Downloading https://files.pythonhosted.org/packages/fa/37/45185cb5abbcb30d7257104c434fe0b07e5a195a6847506c074527aa599ec/Click-7.0-py2.py3-none-any.whl (81kB)
Collecting Jinja2==2.10.1 (from flask->-r requirements.txt (line 1))
  Downloading https://files.pythonhosted.org/packages/1d/e7/fd8b501e7adfe492a433deb7b9d833d39ca74916fa8bc63dd1a4947a671/Jinja2-2.10.1-py2.py3-none-any.whl (124kB)
Collecting itsdangerous==0.24 (from flask->-r requirements.txt (line 1))
  Downloading https://files.pythonhosted.org/packages/76/ae/44b03b253d6fadef317f32c24d100b3b35c2239807046a4c953c7b89fa49e/itsdangerous-1.1.0-py2.py3-none-any.whl
Collecting pvtz (from flask_restful->-r requirements.txt (line 2))
```

# Creating a Runtime from a DockerFile (2/2)

## Dockerfile

```
1 # app.py - a minimal flask api using flask_restful
2 from flask import Flask
3 from flask_restful import Resource, Api
4 app = Flask(__name__)
5 api = Api(app)
6 You, 5 minutes ago | 1 author (You)
7 class HelloWorld(Resource):
8     def get(self):
9         return {'hello': 'world'}
10
11 api.add_resource(HelloWorld, '/')
12 if __name__ == '__main__':
13     app.run(debug=True, host='0.0.0.0') You, 5 minut
```

## Dockerfile

```
You, a few seconds ago | 1 author (You)
1 # Dockerfile - this is a comment. Delete me if you want.
2 FROM python
3 COPY . /app
4 WORKDIR /app
5 RUN pip install -r requirements.txt
6 ENTRYPOINT ["python"]
7 CMD ["app.py"]
8 |
```

## Docker run

```
rmuresano@rmuresano-pc:~/git/bdaasmicroservices/docker/runtimes/helloExample$ docker run --name ronal -d -p 5000:5000 my_docker_flask:latest
eb03c5f48fc4a676dfbc94ed98c30f2593feee515449b11de09dc82c100db534
rmuresano@rmuresano-pc:~/git/bdaasmicroservices/docker/runtimes/helloExample$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
eb03c5f48fc4        my_docker_flask:latest   "python app.py"   3 seconds ago      Up 2 seconds          0.0.0.0:5000->5000/tcp   ronal
```

# DockerFile Instruction

**FROM:** Normally it is the first instruction in the Dockerfile, and it identifies from which image we will start to build the image.

```
1 FROM <image> [AS <name>]
2 FROM <image>[:<tag>] [AS <name>]
3 FROM <image>[@<id-image>] [AS <name>]
```

**ENV:** It is used to declare variables which will be assigned in the building process. It can be used in the following instructions: ADD, COPY, ENV, EXPOSE, FROM, LABEL, USER, VOLUME, WORDIR, etc.

```
1 ENV VAR VALUE
2 ENV VAR=VALUE [VAR2=VALUE]
```

**ARGS:** available during the build of a Docker image (RUN etc), not after the image is created and containers are started from it.

```
1 ARG VERSION=latest
2 FROM busybox:$VERSION
3 ARG VERSION
4 RUN echo $VERSION > image_version
```

# DockerFile Instruction

**RUN:** It executes one or more instruction in the container build. It creates a new layer in case that execution is successfully and this layer will be used by the next Dockerfile instruction.

```
1  FROM ubuntu
2  RUN instruction
3  RUN ["exec file", "param1", "param2", ...]
4  RUN apt-get -y update && apt-get install -y bzip2 build-essential libssl1.0.0 libpam0g openssl \
5      && apt-get install -y autotools-dev libssl1.0-dev libpam0g-dev zlib1g-dev debhelper dh-autoreconf \
6      && apt-get install -y openssh-server sshpass whois net-tools bc nano && \
7      rm -rf /var/lib/apt/lists/*|
```

**EXPOSE:** It is used to inform in which port the container will listen while it is executing. By default, it is TCP, but it can be defined the TCP o UDP

```
9   EXPOSE 8888
10  EXPOSE 8088/UDP|
```

# DockerFile Instruction

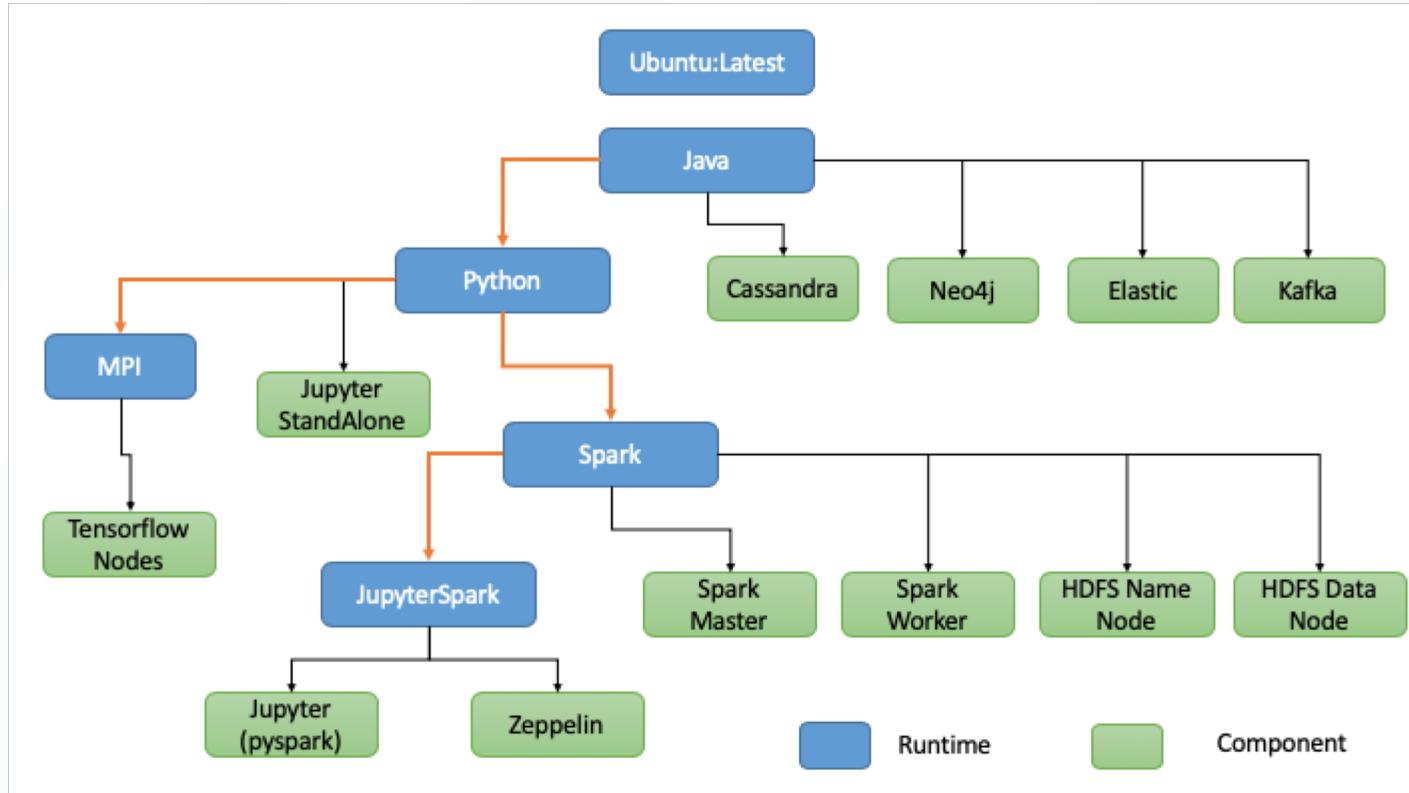
**ENTRYPOINT:** It is used to set up a container as an executable. The instruction indicates which program must be started when the container is started. In case that this instruction appears many times it will be considered the last defined in the Dockerfile.

```
1 ENTRYPOINT ./start-up-config.sh param1 param2
```

**USER:** This instruction establishes the user name (UID) and the group (GID) which must be used when the container processes the following instruction RUN, CMD, ENTRYPOINT.

```
1 RUN cd /src; npm install
2
3 RUN useradd -c 'Node.js user' -m -d /home/node -s /bin/bash node
4 RUN chown -R node.node /src
5 USER node
6 ENV HOME /home/node
7
8 EXPOSE 8080
9 CMD ["node", "/src/index.js"]
```

# DockerFile Runtime Image Tree



# DockerFile Creating the first image (Java)

```
1  FROM ubuntu:latest
2  WORKDIR /home/legacy
3  COPY packages/ ./
4
5  RUN apt-get -y update
6  RUN apt-get install -y bzip2
7  RUN apt-get install -y build-essential
8  RUN apt-get install -y libssl1.0.0
9  RUN apt-get install -y libpam0g
10 RUN apt-get install -y openssl
11 RUN apt-get install -y autotools-dev libssl1.0-dev libpam0g-dev \
12     zlib1g-dev debhelper dh-autoreconf
13 RUN apt-get -y update && \
14     apt-get install -y openssh-server sshpass whois net-tools bc nano && \
15     rm -rf /var/lib/apt/lists/*
16
17 WORKDIR /home/legacy/shellinabox
18 RUN mkdir -p shellinaboxBin && autoreconf -i \
19     && ./configure --prefix=/home/legacy/shellinabox/shellinaboxBin/ \
20     && make && make install
21 You, a few seconds ago • Uncommitted changes
22 ENV PATH="$PATH:/home/legacy/jre/bin"
23
24 WORKDIR /home/legacy
25
```

# DockerFile Creating the first image (Java)

## Example 4: Optimizing the image

[Important] To use a tool for exploring a docker image, layer contents, and discovering ways to shrink your Docker image size.

In our case, we will use Docker Dive: <https://github.com/wagoodman/dive>

Installation Procedure:

### Ubuntu/Debian

```
wget https://github.com/wagoodman/dive/releases/download/v0.7.2/dive_0.7.2_linux_amd64.deb  
sudo apt install ./dive_0.7.2_linux_amd64.deb
```

### Mac

```
brew tap wagoodman/dive  
brew install dive
```

# DockerFile Creating the first image (Java)

Analyzing the image with dive tool

Layers			Current Layer Contents			
Cmp	Size	Command	Permission	UID:GID	Size	Filetree
70 MB	FROM sha256:02571d03		drwxr-xr-x	0:0	4.8 MB	bin
745 B	set -xe && echo '#!/bin/sh' > /usr/sbin/policy-rc.d && echo 'exit 101' >> /usr/sbin/policy-rc.d		-rwxr-Xr-X	0:0	1.1 MB	bash
7 B	mkdir -p /run/systemd && echo 'docker' > /run/systemd/container		-rwxr-Xr-X	0:0	35 kB	bunzip2
0 B	#(nop) WORKDIR /home/legacy		-rwxr-Xr-X	0:0	0 B	bzcat -> bin/bunzip2
241 MB	COPY dir:982f344956926a777e951468f98592ef8832c68392985098e327b15867a96e85 in ./		-rwxrwxrwx	0:0	0 B	bzcmp -> bzdiff
27 MB	apt-get -y update		-rwxr-Xr-X	0:0	2.1 kB	bzdiff
972 kB	apt-get install -y bzip2		-rwxrwxrwx	0:0	0 B	bzegrep -> bzgrep
208 MB	apt-get install -y build-essential		-rwxr-Xr-X	0:0	4.9 kB	bzexe
4.9 MB	apt-get install -y libssl1.0.0		-rwxrwxrwx	0:0	0 B	bzfgrep -> bzgrep
0 B	apt-get install -y libpam0g		-rwxr-Xr-X	0:0	3.6 kB	bzgrep
1.9 MB	apt-get install -y openssl		-rwxr-Xr-X	0:0	0 B	bzip2 -> bin/bunzip2
80 MB	apt-get install -y autotools-dev libssl1.0-dev libpam0g-dev zlib1g-dev debhelper dh-autoreconf		-rwxr-Xr-X	0:0	14 kB	bzip2recover
58 MB	apt-get -y update && apt-get install -y openssh-server sshpass whois net-tools bc nano && rm		-rwxrwxrwx	0:0	0 B	bzless -> bzmore
10 MB	mkdir -p shellinaboxBin && autoreconf -i && ./configure --prefix=/home/legacy/shellinabox/shellinaboxB		-rwxr-Xr-X	0:0	1.3 kB	bzmore
			-rwxr-Xr-X	0:0	35 kB	cat
			-rwxr-Xr-X	0:0	64 kB	chgrp
			-rwxr-Xr-X	0:0	60 kB	chmod
			-rwxr-Xr-X	0:0	68 kB	chown
			-rwxr-Xr-X	0:0	142 kB	cp
			-rwxr-Xr-X	0:0	121 kB	dash
			-rwxr-Xr-X	0:0	101 kB	date
			-rwxr-Xr-X	0:0	76 kB	dd
			-rwxr-Xr-X	0:0	85 kB	df
			-rwxr-Xr-X	0:0	134 kB	dir
			-rwxr-Xr-X	0:0	72 kB	dmesg
			-rwxrwxrwx	0:0	0 B	dnsdomainname -> hostname
			-rwxrwxrwx	0:0	0 B	domainname -> hostname
			-rwxr-Xr-X	0:0	35 kB	echo
			-rwxr-Xr-X	0:0	28 B	egrep
			-rwxr-Xr-X	0:0	31 kB	false

# DockerFile Creating the first image (Java)

## Example 4: Optimizing the image

```
1 Muresano, a day ago | 1 author (Muresano)
2 FROM ubuntu:latest
3 WORKDIR /home/legacy
4 COPY packages/ .
5 RUN apt-get -y update && apt-get install -y bzip2 build-essential libssl1.0.0 libpam0g openssl \
6     && apt-get install -y autotools-dev libssl1.0-dev libpam0g-dev zlib1g-dev debhelper dh-autoreconf \
7     && apt-get install -y openssh-server sshpass whois net-tools bc nano && \
8     rm -rf /var/lib/apt/lists/*
9 WORKDIR /home/legacy/shellinabox
10 RUN mkdir -p shellinaboxBin && autoreconf -i && ./configure --prefix=/home/legacy/shellinabox/shellinaboxBin/ && make && make install
11 ENV PATH="$PATH:/home/legacy/jre/bin"
12
```

Improvement in size around 4.85%

```
rmuresano@rmuresano-pc:~/git/bdaasmicroservices/docker/runtimes/java-optimization-1$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
rmuresano/bdaasjava    0_0_3   edfb976e9539   52 seconds ago  667MB
rmuresano/bdaasjava    0_0_2   01c63bc738a0   8 minutes ago  701MB
```

# DockerFile Creating the first image (Java)

Analyzing the optimized image with dive tool

[● Layers]			[Current Layer Contents]			
Cmp	Size	Command	Permission	UID:GID	Size	Filetree
1	70 MB	FROM sha256:02571d03	-rwxr-Xr-x	0:0	4.8 MB	bash
745 B	set -xe && echo '#!/bin/sh' > /usr/sbin/policy-rc.d && echo 'exit 101' >> /usr/sbin/policy-rc.d &&		-rwxr-Xr-x	0:0	1.1 MB	bunzip2
7 B	mkdir -p /run/systemd && echo 'docker' > /run/systemd/container		-rwxr-Xr-x	0:0	35 kB	bzcat → bin/bunzip2
0 B	#(nop) WORKDIR /home/Legacy		-rwxr-Xr-x	0:0	0 B	bzcmp → bzdiff
241 MB	#(nop) COPY dir:982f344956926a777e951468f98592ef8832c68392985098e327b15867a09e85 in ./		-rwxrwxrwx	0:0	0 B	bzdiff
346 MB	apt-get -y update && apt-get install -y bzip2 build-essential libssl1.0.0 libpam0g openssl && apt-		-rwxr-Xr-x	0:0	2.1 kB	bzegrep → bzgrep
10 MB	10 MB mkdir -p shellinaboxBin && autoreconf -i && ./configure --prefix=/home/legacy/shellinabox/shellinaboxB		-rwxrwxrwx	0:0	0 B	bzexe
			-rwxr-Xr-x	0:0	4.9 kB	bzfgrep → bzgrep
			-rwxrwxrwx	0:0	0 B	bzgrep
			-rwxr-Xr-x	0:0	3.6 kB	bzip2 → bin/bunzip2
			-rwxr-Xr-x	0:0	0 B	bzip2recover
			-rwxr-Xr-x	0:0	14 kB	bzless → bzmore
			-rwxrwxrwx	0:0	0 B	bzmore
			-rwxr-Xr-x	0:0	1.3 kB	cat
			-rwxr-Xr-x	0:0	35 kB	chgrp
			-rwxr-Xr-x	0:0	64 kB	chmod
			-rwxr-Xr-x	0:0	60 kB	chown
			-rwxr-Xr-x	0:0	68 kB	cp
			-rwxr-Xr-x	0:0	142 kB	dash
			-rwxr-Xr-x	0:0	121 kB	date
			-rwxr-Xr-x	0:0	101 kB	dd
			-rwxr-Xr-x	0:0	76 kB	df
			-rwxr-Xr-x	0:0	85 kB	dir
			-rwxr-Xr-x	0:0	134 kB	dmesg
			-rwxr-Xr-x	0:0	72 kB	dnsdomainname → hostname
			-rwxrwxrwx	0:0	0 B	domainname → hostname
			-rwxrwxrwx	0:0	0 B	echo
			-rwxr-Xr-x	0:0	35 kB	egrep
			-rwxr-Xr-x	0:0	28 B	false
			-rwxr-Xr-x	0:0	31 kB	

^C Quit | Tab Switch view | ^F Filter | ^I Show layer changes | ^A Show aggregated changes |

# DockerFile Creating the first image (Java)

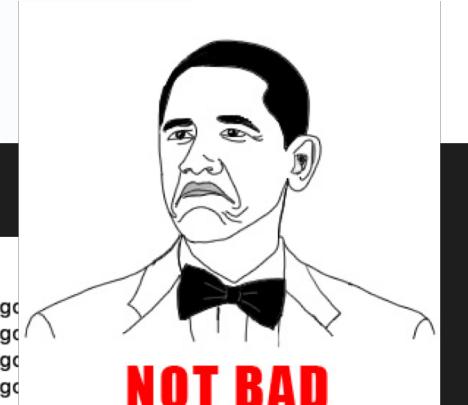
To avoid using different Dockerfiles, it was add the functionality to define diverse steps in the same file. Each step start with the instruction from.

Once a step is built, we can copy files from other layers using the **COPY** command and the parameter flag --from

Example:

Improvement in size around 54 %

```
Muresano, a day ago | 1 author (Muresano)
1 FROM ubuntu:latest as builder
2 WORKDIR /home/legacy
3 COPY packages/ .'
MacBook-Pro-de-Ronal:radius3 rmuresano$ docker images
4 RUN apt-get -y update
REPOSITORY          TAG      IMAGE ID      CREATED
5 && apt-get install rmuresano/bdaascassandra    0_0_4      23685696b08f  23 hours ago
6 && apt-get install rmuresano/bdaasneo4j        0_0_4      ed8b92379679  25 hours ago
7 rm -rf /var/..
rmuresano/bdaassparkworker   0_0_4      d1ff45a0f6ae  29 hours ago
8 WORKDIR /home/legacy
rmuresano/bdaassparkmaster  0_0_4      16f68f5371c0  29 hours ago
9 RUN mkdir -p shell
rmuresano/bdaasjupyter     0_0_4      5d359d0e0935  29 hours ago
10 rmuresano/bdaasspark
rmuresano/bdaaspark        0_0_4      3cfdece2be29  30 hours ago
11 FROM ubuntu:latest
rmuresano/bdaaspthon       0_0_4      151df2f5b132  30 hours ago
12 WORKDIR /home/legacy
rmuresano/bdaasjava        0_0_4      fa78e638247f  30 hours ago
13 COPY --from=builder:0_0_4 rmuresano/bdaasjava    0_0_3      cf4cb6d87e93  30 hours ago
14 COPY --from=builder:0_0_2 rmuresano/bdaasjava    0_0_2      87ffd9b93379  30 hours ago
15 ENV PATH="$PATH:$PWD/jupyter/datascience-notebook"
ubuntu                  latest    7698f282e524  7 weeks ago   69.9MB
jupyter/datascience-notebook latest    029fd3e52059  8 weeks ago   5.49GB
16
17
```



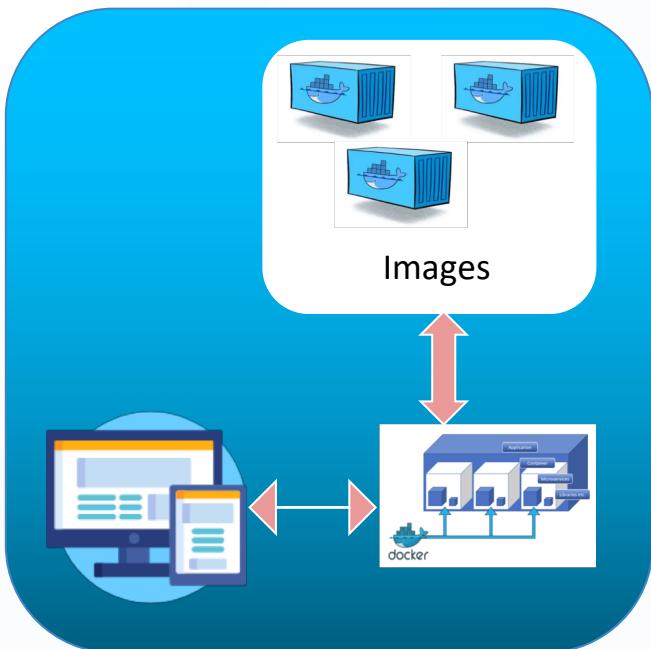
# DockerFile Creating the Second image (Python)

Exercise 1: Creating the first image. Time Approximate (25~30 Min)

Steps:

- 1: Download the miniconda packages <https://docs.conda.io/en/latest/miniconda.html>
- 2: Create an image from the optimized image rmuresano/bdaasjava:0\_0\_4
- 3: Copy the installer to the image
- 4: Install the
- 4: Update the conda and pip packages
- 5: Remove all additional packages
- 6: Establish and ENV variable with the miniconda/bin Path.

# Docker Registry (Runtime for BDaaS)



- ❖ Docker Trusted Registry

- ❖ Docker Hub
- ❖ Docker Trusted on Cloud
- ❖ Azure container Registry
- ❖ AWS container Registry (ECS -ECR)
- ❖ Google Container Registry.
- ❖ Other Cloud

On-premise

Public Clouds

# Docker Hub Registry (Creating an account)

Exercise 2: <https://hub.docker.com/> Time 5~10 minutes



### Docker Identification

In order to get you started, let us get you a Docker ID.  
Already have an account? [Sign In](#)

Enter a Docker ID

Password

Email

I agree to Docker's [Terms of Service](#).

I agree to Docker's [Privacy Policy](#) and [Data Processing Terms](#).

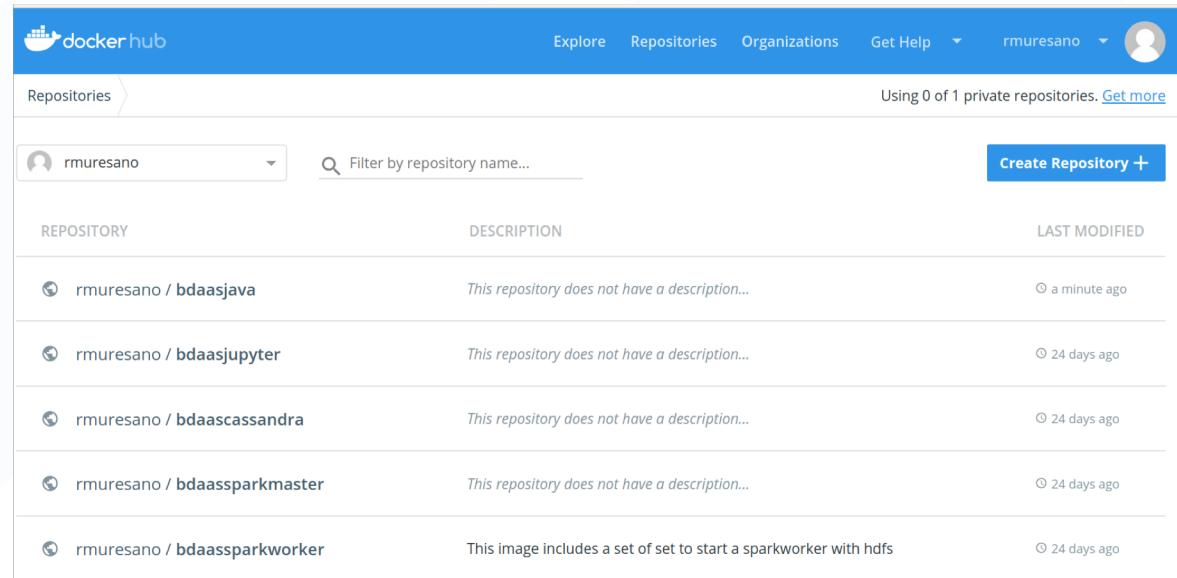
(Optional) I would like to receive email updates from Docker, including its various services and products.

No soy un robot   
reCAPTCHA  
Privacy + Conditions

**Continue**

# Push image to docker Hub

```
rmuresano@rmuresano-PC:~/git/bdaasmicroservices/docker/runtimes/java-optimization-2$ docker push rmuresano/bdaasjava:0_0_4
The push refers to a repository [docker.io/rmuresano/bdaasjava]
7f63c72cccd1b: Pushed
07da8bc77580: Pushing [=====] 239.3MB
cd9a495985a2: Pushed
8d267010480f: Pushed
270f934787ed: Pushed
02571d034293: Pushed
```



The screenshot shows the Docker Hub user interface for the user 'rmuresano'. The top navigation bar includes links for Explore, Repositories, Organizations, Get Help, and a user profile icon. A message indicates 'Using 0 of 1 private repositories.' with a 'Get more' link. Below this, a dropdown menu shows 'rmuresano' and a search bar with the placeholder 'Filter by repository name...'. A 'Create Repository +' button is visible. The main content area displays a table of repositories:

REPOSITORY	DESCRIPTION	LAST MODIFIED
rmuresano / bdaasjava	This repository does not have a description...	⌚ a minute ago
rmuresano / bdaasjupyter	This repository does not have a description...	⌚ 24 days ago
rmuresano / bdaascassandra	This repository does not have a description...	⌚ 24 days ago
rmuresano / bdaassparkmaster	This repository does not have a description...	⌚ 24 days ago
rmuresano / bdaassparkworker	This image includes a set of set to start a sparkworker with hdfs	⌚ 24 days ago

# Deploy a registry server

- 1) Use a command like the following to start the registry container:

```
$ docker run -d -p 5000:5000 --restart=always --name registry registry:2
```

- 2) Pull the ubuntu:16.04 image from Docker Hub.

```
$ docker pull ubuntu:16.04
```

- 3) Tag the image as localhost:5000/my-ubuntu

```
$ docker tag ubuntu:16.04 localhost:5000/my-ubuntu
```

- 4) Tag the image as localhost:5000/my-ubuntu

```
$ docker push localhost:5000/my-ubuntu
```

# Compile Components and Runtimes tree

Exercise 3: 15~20 minutes to create all environment

Objective: Learn the procedure to compile and create big data components.

Steps:

Runtime

- 1) Git clone <git@github.com>:mureron/bdaasmicroservices.git bdaas.
- 2) Enter in the docker/runtime/packages directory.
- 3) Run the script ./DownloadFiles.sh (wait until all software and dependencies are downloaded).
- 4) Exit the docker/runtime/packages directory and enter the docker/runtime
- 5) Execute the script build-runtime.sh. (Take a coffee)

Components

- 1) Enter in the folder docker/components
- 2) Execute the script ./compileComponents.sh
- 3) Now we have all the runtimes and components done :)

# Course Roadmap

**Orchestration for developing Docker and Docker-Compose.**

## Big Data Challenges

An overview applied to the SMEs.  
Orchestration Philosophy

## Docker Swarm

Orchestration in a real private cluster (Deploying and Creating services)

## Kubernetes

Production Orchestration.  
Automating deployment, scaling, and management of containerized applications

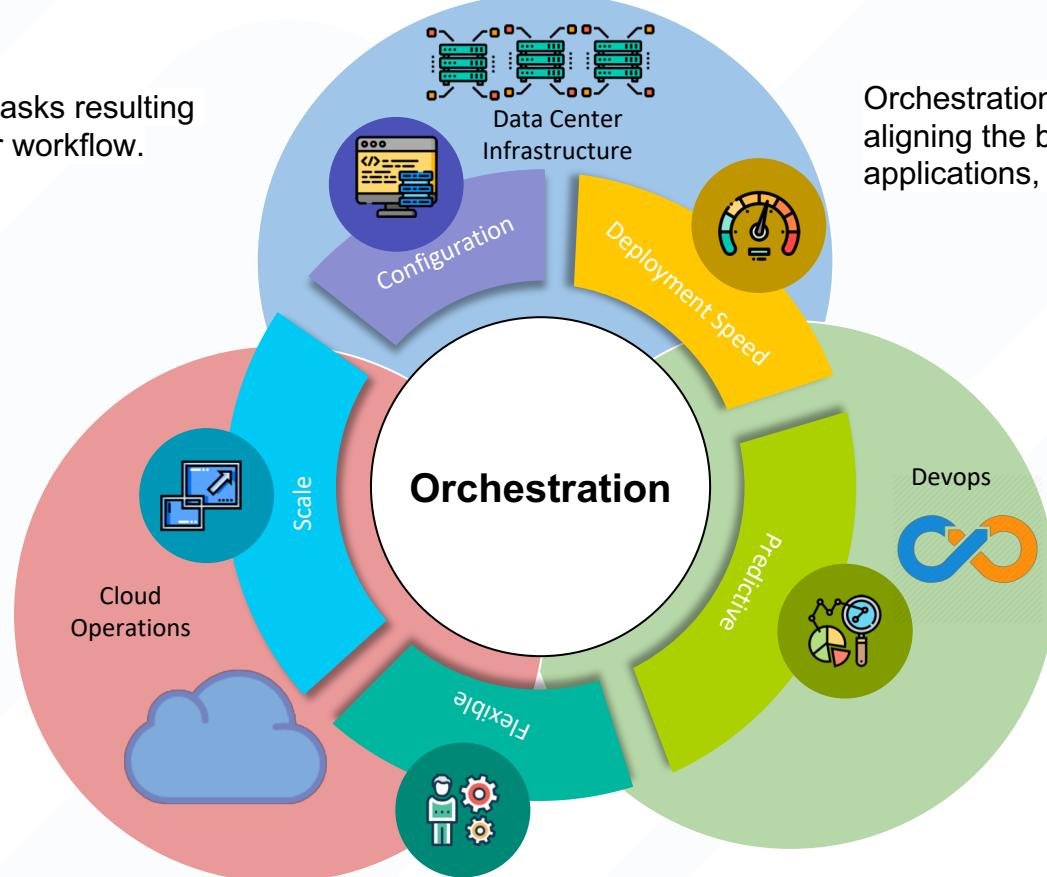
## Cloud Infrastructure

Google Cloud and AWS deployments



# Service Orchestration

Coordination of automated tasks resulting in a consolidated process or workflow.



Orchestration in this sense is about aligning the business request with the applications, data, and infrastructure.

# Docker-Compose (Multi-container Services)

Typical applications have multiple components. For example, Web App.

Docker Compose is a tool for defining and running multi-container Docker application.

with a single command, you create and start all the services from your configuration.

Compose works in all environments: production, staging, development, testing, as well as CI workflows

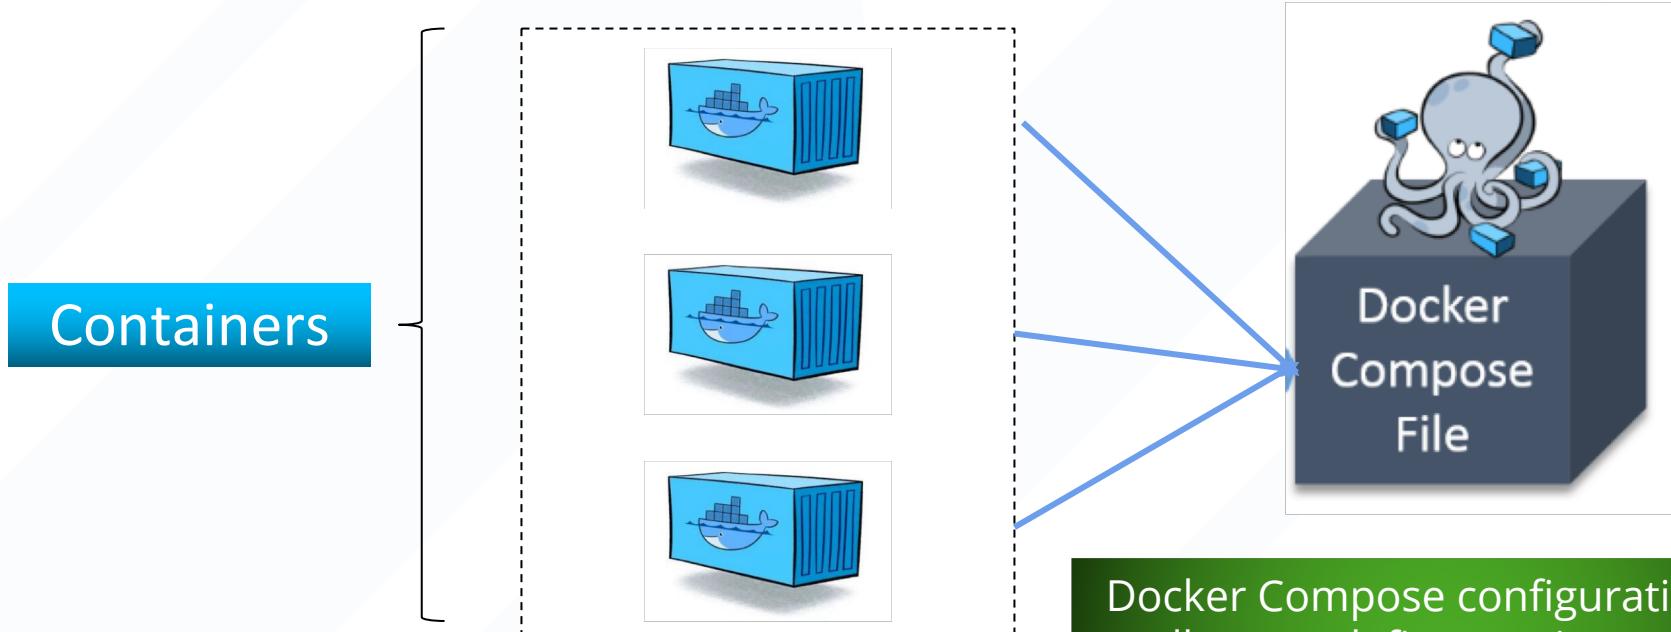
Using Compose is basically a three-step process:

1. Define your app's environment with a Dockerfile so it can be reproduced anywhere.
2. Define the services that make up your app in docker-compose.yml so they can be run together in an isolated environment.
3. Run docker-compose up and Compose starts and runs your entire app.

# Docker-Compose (Multi-container Services)

A first Step from the development to a production environment.

The main idea is to help you to define and build application stacks.



Docker Compose configuration file  
allows to define: environment  
variables, networks or volumes

# Docker-Compose (Multi-container Services)

```
git clone https://github.com/rmuresan/bdaas-docker-compose.git  
cd bdaas-docker-compose  
# Start the JupyterLab service  
$ docker-compose up jupyter
```

```
version: '3'  
services:  
  jupyter:  
    image: rmuresano/bdaasjupyter:0_0_4  
    hostname: jupyter  
    container_name: jupyter  
    environment:  
      - STANDALONE=YES  
      - JUPYTERPORT=8900  
      - HDFS=NO  
    ports:  
      - "8900:8900" # WEB interface JupyterLab  
    volumes:  
      - jupyter-vol:/data/jupyter/  
  
volumes:  
  jupyter-vol:
```

- ❖ The cluster specification is done in a YAML file, normally called docker-compose.yml

- ❖ It is started using

```
docker-compose up
```

- ❖ The service can be scaled using the option

```
docker-compose up --scale SERVICE=NUM
```

Demo

# Running Docker-Compose (1/5)

Example 5: Start a batch service for Big Data (Jupyter-Spark)

```
MacBook-Pro-de-Ronal:jupyterspark rmuresano$ docker-compose up -d
```

```
Creating sparkmaster ... done
```

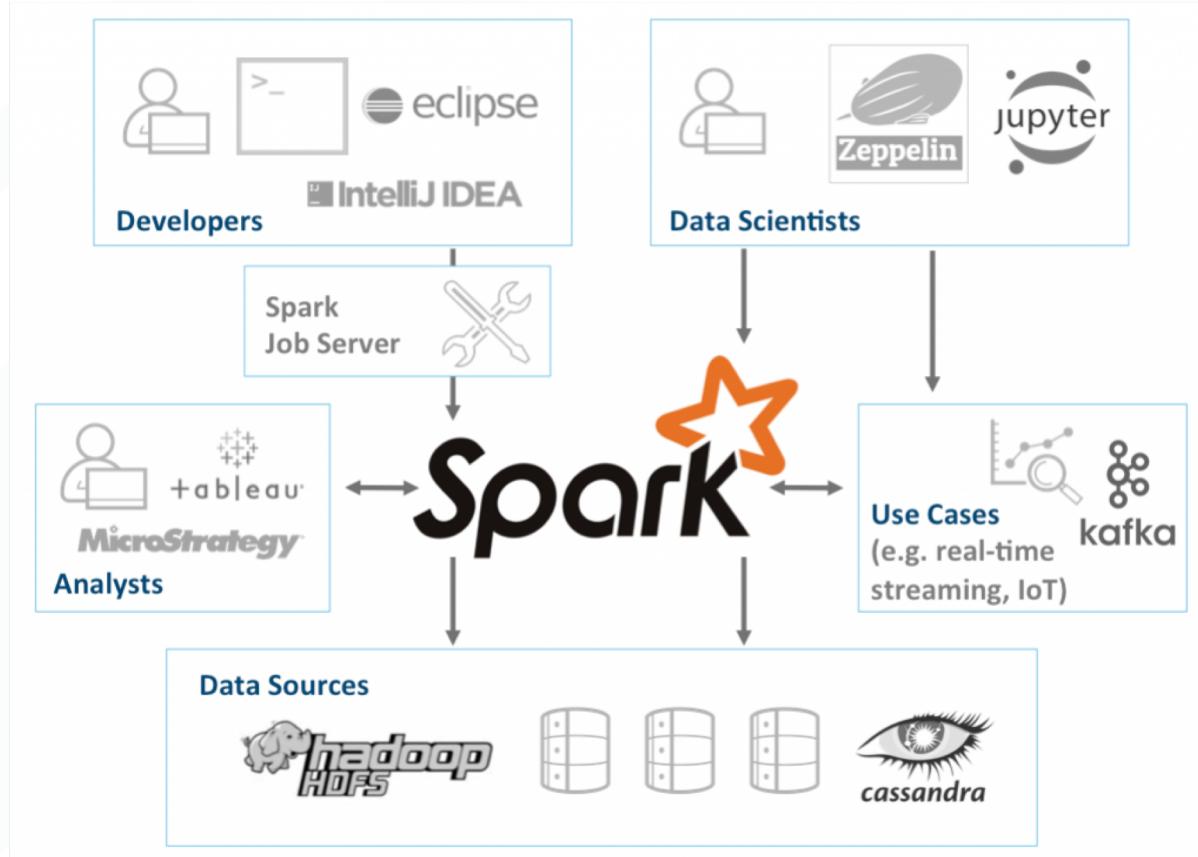
```
Creating jupyter ... done
```

```
Creating jupyterspark_sparkworker_1 ... done
```

```
MacBook-Pro-de-Ronal:jupyterspark rmuresano$ docker-compose ps
```

Name	Command	State	Ports
jupyter	/bin/sh -c ./startup-config.sh	Up	0.0.0.0:8900->8900/tcp
jupyterspark_sparkworker_1	/bin/sh -c ./startup-config.sh	Up	0.0.0.0:8087->8081/tcp
sparkmaster	/bin/sh -c ./startup-config.sh	Up	0.0.0.0:50070->50070/tcp, 0.0.0.0:7077->7077/tcp, 0.0.0.0:8080->8080/tcp, 0.0.0.0:9000->9000/tcp

# Running Docker-Compose (2/5)

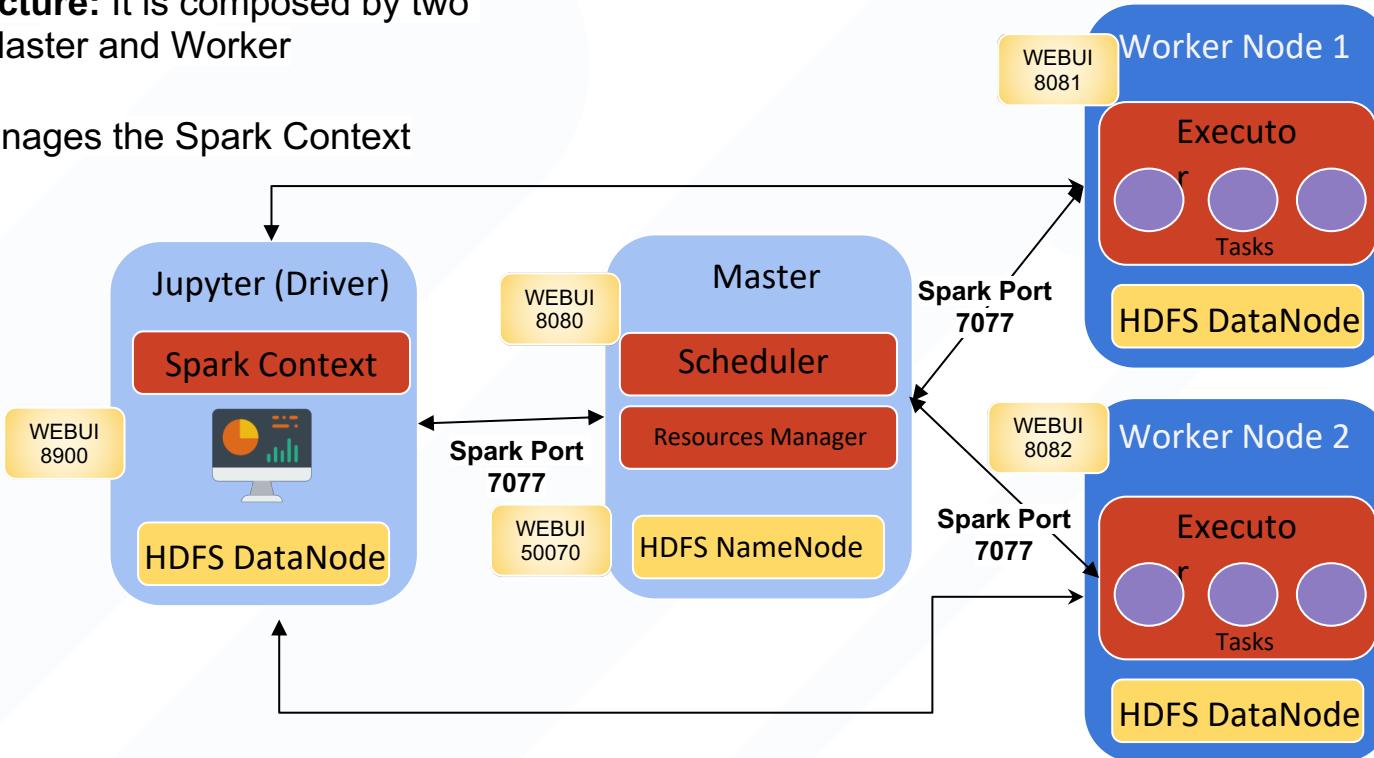
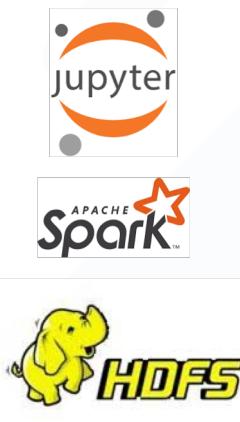


# Running Docker-Compose (3/5)

Analyzing the architecture, we have to split into two parts (Jupyter, Spark and HDFS))

**Spark architecture:** It is composed by two components Master and Worker

**Jupyter:** It manages the Spark Context



# Running Docker-Compose (4/5)

## Jupyter

**Environment:**

```
STANDALONE=NO  
JUPYTER PORT=8900  
SPARK_MASTER_HOSTNAME=sparkmaster  
SPARK_MASTER_PORT=7077  
HDFS_NAMENODE_HOSTNAME=sparkmaster  
HDFS_NAMENODE_METADATA_PORT=9000  
HDFS_REPLICATION_FACTOR=1  
HDFS_DATANODE_WEBUI_PORT=50075  
SPARK_WORKER_WEBUI_PORT=8081  
HDFS_NAMENODE_WEBUI_PORT=50070  
HDFS=YES
```

**Ports to Expose:**

```
8900:8900
```

**Volume:**

```
jupyter-vol:/data/jupyter
```

**Depends on:**

```
Spark Master
```

## Spark Master

**Environment:**

```
SPARK_MASTER_PORT=7077  
SPARK_MASTER_WEBUI=8080  
HDFS_CLUSTER_NAME=HDFS_TEST  
HDFS_NAMENODE_METADATA_PORT=9000  
HDFS_REPLICATION_FACTOR=1  
HDFS_NAMENODE_WEBUI_PORT=50070  
HDFS=YES
```

**Ports to Expose:**

```
8080:8080  
7077:7077  
9000:9000  
50070:50070
```

**Volume:**

```
hdfs-name-node:/hdfs/namenode
```

## Spark Worker

**Environment:**

```
SPARK_MASTER_HOSTNAME=sparkmaster  
SPARK_MASTER_PORT=7077  
HDFS_NAMENODE_HOSTNAME=sparkmaster  
HDFS_NAMENODE_METADATA_PORT=9000  
HDFS_REPLICATION_FACTOR=1  
HDFS_DATANODE_WEBUI_PORT=50075  
SPARK_WORKER_WEBUI_PORT=8081  
HDFS_NAMENODE_WEBUI_PORT=50070  
HDFS=YES
```

**Ports to Expose:**

```
8081-8090:8081
```

**Volume:**

```
hdfs-data-node:/hdfs/datanode
```

**Depends on:**

```
Spark Master
```

# Running Docker-Compose (5/5)

Spark Master

```
# Set up NameNode and Spark Master
sparkmaster:
  image: rmuresano/bdaassparkmaster:0_0_4
  hostname: sparkmaster
  container_name: sparkmaster
  environment:
    - SPARK_MASTER_PORT=7077
    - SPARK_MASTER_WEBUI_PORT=8080
    - HDFS_NAMENODE_METADATA_PORT=9000
    - HDFS_CLUSTER_NAME=HDFS_TEST
    - HDFS_NAMENODE_WEBUI_PORT=50070
    - HDFS_REPLICATION_FACTOR=1
    - HDFS=YES
  ports:
    - "8080:8080" #Web Interface Master
    - "7077:7077" #Master Port
    - "9000:9000" # HDFS PORT
    - "50070:50070" # WEB NameNode Int
  volumes:
    - hdfs-name-node:/hdfsdata/nameNode/
```

Spark Worker

```
# Setup the Worker Node 1.
sparkworker:
  image: rmuresano/bdaassparkworker:0_0_4
  hostname: sparkworker
  depends_on:
    - sparkmaster
  links:
    - sparkmaster
  environment:
    - SPARK_MASTER_HOSTNAME=sparkmaster
    - SPARK_MASTER_PORT=7077
    - HDFS_NAMENODE_HOSTNAME=sparkmaster
    - HDFS_NAMENODE_METADATA_PORT=9000
    - HDFS_REPLICATION_FACTOR=1
    - HDFS_DATANODE_WEBUI_PORT=50075
    - SPARK_WORKER_WEBUI_PORT=8081
    - HDFS_NAMENODE_WEBUI_PORT=50070
    - HDFS=YES
  ports:
    - "8081-8090:8081"
```

Jupyter

```
jupyter:
  image: rmuresano/bdaasjupyter:0_0_4
  hostname: jupyter
  container_name: jupyter
  depends_on:
    - sparkmaster
  links:
    - sparkmaster
  environment:
    - STANDALONE=NO
    - JUPYTERPORT=8900
    - HDFS=YES
    - HDFS_NAMENODE_HOSTNAME=sparkmaster
    - HDFS_DATANODE_WEBUI_PORT=50075
    - HDFS_NAMENODE_METADATA_PORT=9000
    - HDFS_REPLICATION_FACTOR=1
    - SPARK_MASTER_HOSTNAME=sparkmaster
    - SPARK_MASTER_PORT=7077
  ports:
    - "8900:8900" # WEB interface JupyterLab
  volumes:
    - jupyter-vol:/data/jupyter/
```

Run the example using: Docker-compose up --scale sparkworker=2

# Running Docker-Compose (5/5)

The screenshot shows a web browser window with two tabs open:

- localhost:8900/lab?**: A Jupyter Notebook interface. On the left, a file tree lists several Python notebooks (ipynb files) with their names and last modified dates. On the right, there are three main sections: "Notebook" (Python 3), "Console" (Python 3), and "Other".
- localhost:8080**: The Spark Master UI. It displays the URL `spark://sparkmaster:7077`. Key metrics shown include: Alive Workers: 1, Cores in use: 2 Total, 0 Used, Memory in use: 1024.0 MB Total, 0.0 B Used, Applications: 0 Running, 0 Completed, Drivers: 0 Running, 0 Completed, and Status: ALIVE.

**Workers (1)**

Worker Id	Address	State	Cores	Memory
worker-20190709115807-172.25.0.3-44823	172.25.0.3:44823	ALIVE	2 (0 Used)	1024.0 MB (0.0 B Used)

**Running Applications (0)**

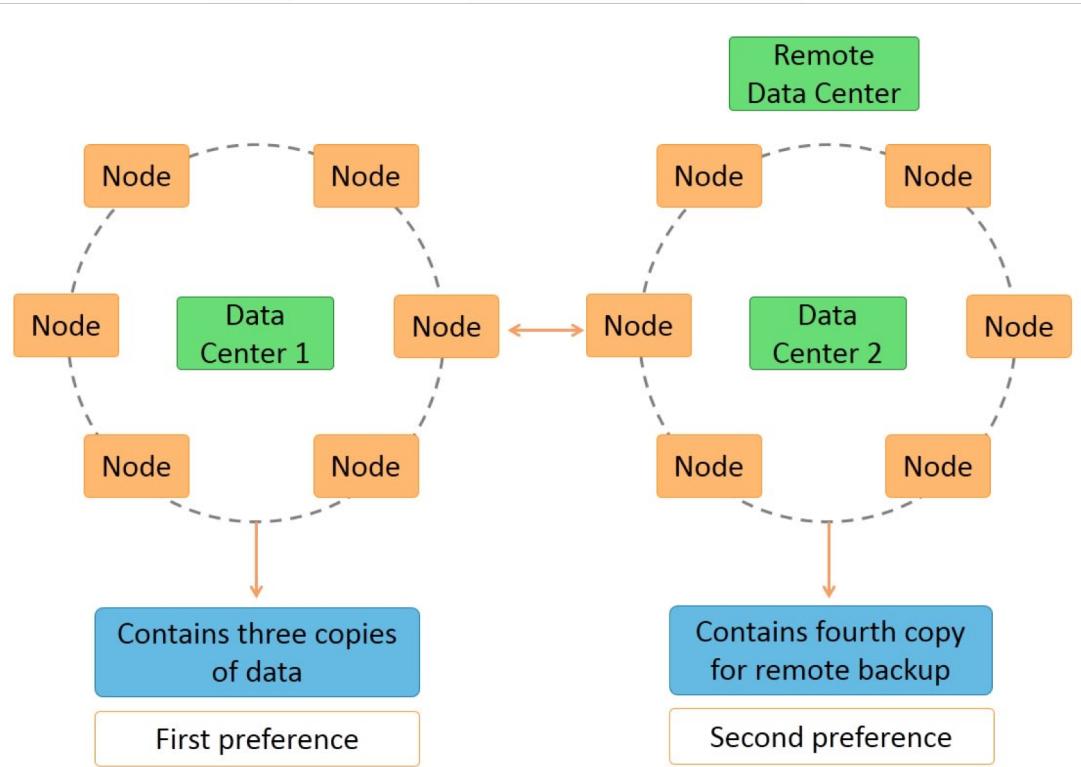
Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

**Completed Applications (0)**

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

# Running another example

Exercise 4: Time 20 minutes (How to create a compose for Jupyter-Cassandra Cluster )



```
cassandra-seed:  
image: rmuresano/bdaascassandra:0_0_4  
environment:  
- CLUSTER_NAME=seminario  
- STORAGE_PORT=7000  
- NATIVE_TRANSPORT_PORT=9042  
- RPC_PORT=9160  
- CASSANDRA_SEEDS=172.28.1.2,172.28.1.3,172.28.1.4  
volumes:  
- cassandra-volume-seed:/data/cassandra  
networks:  
cassandra:  
| ipv4_address: 172.28.1.2
```

## Big Data as a Service: A new approach for the SMEs

Muchas gracias por su atención

Ronal Muresano

[rmuresano@gmail.com](mailto:rmuresano@gmail.com)

[rmuresano@iti.es](mailto:rmuresano@iti.es)

